

Answersheet

Part-1

① Worst-case time complexity of finding the maximum element in a min-heap $\Rightarrow O(n)$ \rightarrow leaf nodes.

② Is $\log_e(n) = \Theta(\log_2(n))$? Yes/No

Answer: Yes, $\log_e(n) = \frac{\log_2(n)}{\log_2(e)} = \frac{1}{\log_2(e)} \cdot \log_2(n)$

$$\therefore \log_e(n) = c \cdot \log_2(n)$$

$$\left\{ \text{where } c = \frac{1}{\log_2(e)} \right\} \Rightarrow \boxed{\log_e(n) = \Theta(\log_2(n))}$$

by definition of Θ -notation,

$$f(n) = \Theta(g(n)) \Rightarrow \exists c_1, c_2 > 0 \text{ & } n_0 \text{ s.t. }$$

$$\Rightarrow c_1 g(n) \leq f(n) \leq c_2 g(n) \text{ for } n \geq n_0.$$

③ If A is an-element array where each element is b-bit integer, the worst-case complexity of sorting A using the best algorithm depends on the constraints:

* Comparison-based Sorting : Merge, Quick or Heap sort

→ Worst-case : $O(n \lg n)$.

→ This is because comparison-based sorting is bounded by $\Omega(n \lg n)$ in the worst case, regardless of the bit-size b of integers.

* Non-comparison-based sorting

If we exploit the facts that integers are b -bit values.
thus range $[0, 2^b - 1]$.

→ Radix Sort : $O(nb)$

→ Counting Sort : $O(n+k)$, where $k=2^b$.

Hence, best case is Radix Sort for b -bits i.e. $O(nb)$.

④ Dijkstra's Algorithm is a greedy algorithm

⑤ Yes, the problem of determining whether there exist a path between 2 vertices s & t in an undirected graph $G = (V, E)$ is in NP, but it is in P.

→ NP membership ↗ A soln can be verified in polynomial time given a "certificate".

The verification process occurs in polynomial time.

For the given problem:

→ The certificate is the sequence of vertices (a path) that connects s to t .

→ The verification involves checking whether the given sequence is a valid path in G .

↳ Verifying that each pair of consecutive vertices in the sequence has an edge in G .

↳ Checking that the path starts at s and ends at t .

Ex:- BFS & DFS $\not\propto O(V+E)$.

Part-2

⑥ To solve the problem of finding an index i such that $A[i] = i$ in a sorted array $A[1, 2, \dots, n]$, we can leverage the sorted property of A to design an efficient algo. i.e. Binary search.

Algorithm:

- ① Initialize 2 pointers $low = 1$ & $high = n$.
- ② while $low \leq high$:
 - ① compute the middle index : $mid = \lfloor (low+high)/2 \rfloor$
 - ② check if $A[mid] = mid$, return mid
 - if $A[mid] < mid$, set $low = mid + 1$
 - if $A[mid] > mid$, set $high = mid - 1$

③ If the loop exits without finding $A[i] = i$, return "None".

* Time complexity: $O(n \log n)$.

(7) To prove that at least $n-1$ comparisons are needed to check whether all the elements in an array of n elements are the same, let's use a decision tree argument.

* Understanding the problem:

We have an array $A[1], A[2], \dots, A[n]$ & want to determine if all elements are the same.

The key operation is comparing 2 elements $A[i]$ & $A[j]$, which results in either:

→ Equal($A[i] == A[j]$) or

→ Not equal($A[i] \neq A[j]$)

If any 2 elements are not equal, we can immediately conclude that the array elements are not the same. However, to ensure all n elements are equal, we must systematically check every pair.

Approach 2: Lower bound using decision tree

A decision tree is binary tree representing the sequence of comparisons made & their outcomes (equal or not equal). Each leaf node corresponds to a final decision (all elements are the same or not).

* Proof:

① Base Case: $n=2$, we need exactly 1 comparison $A[1] \neq A[2]$

$$T(n) = n-1$$

② Inductive Step: Assume that for an array of size k , at least $k-1$ comparisons are required. Now, consider an array of size $k+1$:

→ Compare $A[1]$ with $A[2], A[3], \dots, A[k-1]$ sequentially

→ To confirm all elements are equal, we must verify that each $k+1$ elements matches the 1st element.

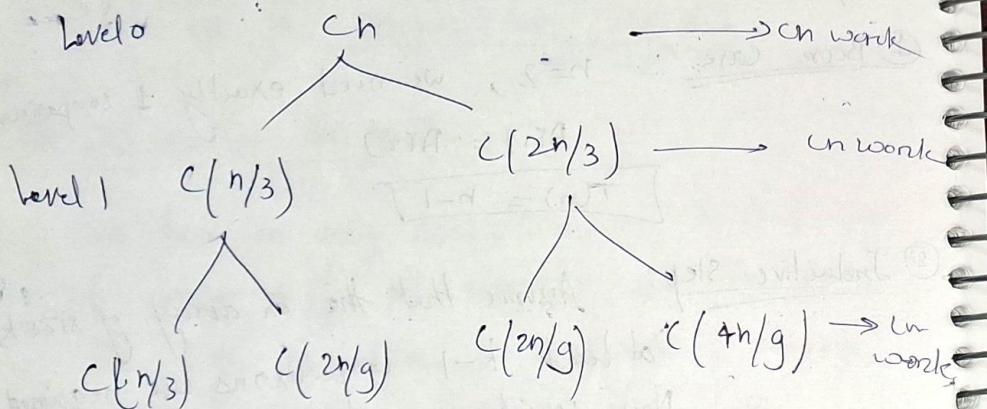
By the inductive hypothesis, verifying the first k elements requires $k-1$ comparisons. Adding the comparison for $(k+1)$ th element requires one additional comparison.

resulting $k+1 - 1 = k$ comparisons.

Thus, the recursion holds for $k+1$.

Hence proved.

(8) $T(n) = T(n/3) + T(2n/3) + cn$



Total cost = $c \cdot \frac{n}{9} + c \cdot \frac{2n}{9} + c \cdot \frac{2n}{9} + c \cdot \frac{4n}{9}$

= cn.

Given: Depth = 3 (Given cn) $\left\{ \begin{array}{l} T(n) = (cost per level) * \\ \quad | \text{No. of levels} \\ \text{if } (cn \cdot \log_3 n) \end{array} \right.$

$\Rightarrow d = \log_3(n).$

Ans =

⑨ Mistake in the Inductive Proof of $T(n) = 2T(n/2) + cn$.

Given claim: The solution incorrectly concludes

$$T(n) = O(n) \text{ using induction.}$$

+ Mistake:

The key mistake lies in the inductive step:

→ The assumption $T(n) = O(n)$ $\forall n < k$ if

→ When expanding $T(k) = 2T(k/2) + \Theta(k)$, they directly replace $T(k/2) = O(k/2)$ resulting in $2 \cdot O(k/2) + \Theta(k) = \Theta(k)$.

This reasoning is "flawed" because it ignores the multiplicative growth of the recurrence.

Specifically: $T(n) = 2T(n/2) + cn \neq O(n)$ dominated by $O(n)$, as the recurrence grows logarithmically in depth.

Correct Sol: Using master theorem, $T(n) = \Theta(n \log n)$

$$a=2, b=2 \text{ and } d=1$$

$$a=b \Rightarrow d=2$$

Case 2: $T(n) = \Theta(n \log \log n)$

$T(n) = \Theta(\log^2 n)$

10 Algorithm to check whether a graph is Bipartite

→ Approach 1: DFS & BFS

→ Approach 2: 2-coloring method.

A bipartite graph can be defined as a graph whose vertices can be divided into 2 disjoint sets such that no 2 vertices in the same set are connected by an edge. This is equivalent to saying that the graph can be 2-colored, meaning that its vertices can be assigned one of 2 colors (i.e. 0 or 1) such that no two adjacent vertices share the same color.

Algorithm:

Step(1): Initialize coloring array:

→ Create an array $\text{color}[]$ of size V (no. of vertices).

→ Initialize all vertices as uncolored i.e.

$$\text{Color}[v] = -1 \quad \forall v \in V.$$

Step(2): Iterate through all vertices:

→ Since the graph may be disconnected, iterate through all vertices.

→ If a vertex v is not colored, start a BFS/DFS from v & assign it the first color (e.g. $\text{color}[v] = 0$)

Step(3): Perform BFS/DFS for 2-coloring.

→ During traversal, for every edge (u, v) :

 → If v not colored \Rightarrow assign opposite color of u .
 i.e. $\text{color}[v] = 1 - \text{color}[u]$.

 → If already colored $\Rightarrow \text{color}[v] = \text{color}[u]$,
 the graph is not bipartite (as adjacent vertices share the same color).

Step(4): Repeat for all components.

→ If BFS/DFS completes without conflicts for all components, the graph is bipartite.

$$\boxed{\text{Time Complexity} = \mathcal{O}(V+E)}$$

2nd Approach (BFS):

Algorithm:

Step ① Initialize an array $\text{color}[]$ with size equal to the No. of vertices. Set all values to -1 (uncolored).

Step ② For each unvisited vertex v :

 ① Assign $\text{color}[v] = 0$ & start a BFS from v .

 ② During BFS, for each edge (u, v) :

 If $\text{color}[v] = -1$, assign ~~color[v] = color[u]~~
 assign $\text{color}[v] = 1 - \text{color}[u]$

 If $\text{color}[v] = \text{color}[u]$,

 the graph is not bipartite.

Step ③ If the BFS completes without conflicts,
the graph is bipartite.

* Time complexity = $O(V+E)$.
 No. of Vertices No. of Edges

11

Algorithm for minimum-weight Spanning tree with negative weights.

Algorithm : Prim's OR Kruskal's Algorithm

Kruskal's Algorithm

① Sort all edges by weight (including $-ve$)

② Initialize a disjoint-set data structure (Union-Find) for the graph vertices.

③ Iterate through the sorted edges:

 If an edge does not form a cycle,
 add it to the MST.

④ Stop when $V-1$ edges are added to the MST.

* Time complexity:

→ Sorting the edges: $O(E \log E)$.

→ Union-Find Operations: $O(E \alpha(V))$, where $\alpha(V)$
 is the inverse Ackermann function.

Hence, $\boxed{TC(n) = O(E \log E)}$ Ans.

12) Decreasing -key in a min-heap involves:

- ① Updating the key values of specific element.
- ② Restoring the heap property by moving the updated element upward (heapsify-up).

* Time complexity:

- In a binary min-heap, the height is $O(\log n)$.
- Thus, the time complexity of decreasing a key is $O(\log n)$.

13) a) Given a directed acyclic graph (DAG), analyze the properties of the adjacency matrix A.

- ① $A^n = 0$, where n is the no. of vertices. This is because no path in a DAG can contain more than $n-1$ edges.

- ② The power of A (e.g. A^k) indicates the no. of paths of length k b/w any pair of vertices.

b) NP-completeness of Clique₁₀₀

determine if Clique₁₀₀, the problem of deciding whether has a clique of size 100, is NP-complete.

Answer: → Clique₁₀₀ is in NP because a certificate (subset of 100 vertices) can be verified in polynomial time.

→ To prove NP-completeness:

→ The general clique problem is NP-complete.

→ Clique₁₀₀ is specific instance of clique with $k=100$, so it is also NP-complete.

13) Reducing L to SAT or SAT to L

To prove the L is NP-hard, should we reduce SAT to L or L to SAT?

Answer: To prove L is NP-hard → Reduce SAT to L.

→ If SAT (a known NP-complete problem) can be reduced to L in polynomial time, the L is at least as hard as SAT.