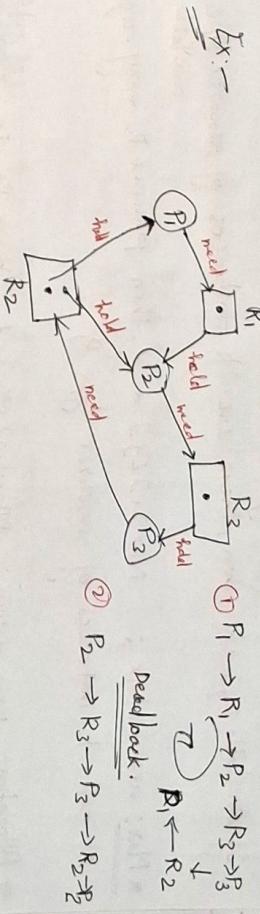


Method to Recover from Deadlock

- ① Process Termination → Abort specific processes
- ② Resource Preemption \hookrightarrow Break the circular dependency
- ③ Detection & Recovery
- ④ Deadlock avoidance.
- ⑤ Deadlock prevention

Absent all processes may complete from a finite no. of resources. A process request resources; if the resource is available at that time a process enters the wait state. Waiting process may never change its state because the resources requested are held by other waiting process. This situation is known as deadlock.

① Deadlock: In a multiprogramming env several processes may compete from a finite no. of resources. A process request resources; if the resource is available at that time a process enters the wait state. Waiting process may never change its state because the resources requested are held by other waiting process. This situation is known as deadlock.



Necessary & Sufficient condition for a deadlock (Characteristics)

* ① Mutual Exclusion: At a time only one process can use the resource. If another process requests that resource, requesting process must wait until the resource has been released.

* ② Hold & Wait: A process must be holding at least one resource & waiting for additional resource that is currently held by other processes.

* ③ No preemption: Resources allocated to a process can't be forcibly taken out from it unless it releases that resource after completing the task.

* ④ Circular wait: A set $\{P_0, P_1, \dots, P_n\}$ of waiting processes must exists such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for the resource that is held by P_2 , ..., P_{n-1} is waiting for resource that is held by P_n & P_n is waiting

for resources that is held by P_i .

Ex:- "Dining Philosophers" in deadlock.

② Banker's Algoithm: Deadlock avoidance algorithm

Let $n \rightarrow$ No. of processes, $m \rightarrow$ No. of resources types.

* Available: Vector of length m .

If $\text{Available}[j] = k$, then there are k instances of resource type R_j available.

* Max: $n \times m$ matrix; $\text{Max}[i][j] = k$, then process P_i may request at most k instances of resource type R_j .

* Allocation: $n \times m$ matrix; $\text{Allocation}[i][j] = k$ then P_i is currently allocated k instances of resources type R_j .

* Need: $n \times m$ matrix; $\text{Need}[i][j] = k$ then P_i may need k more instances of R_j to complete its task.

$\text{Need}[i][j] = \text{Max}[i][j] - \text{Allocation}[i][j]$

* Finish: Boolean values, either TRUE or FALSE.

If $\text{finish}[i] = \text{TRUE}$ & i return safe state else unsafe.

Safety Condition:

Deadlock	Unsafe
----------	--------

Step(2): Find an i such that both:

(a) $\text{Finish}[i] = \text{False}$
 (b) $\text{Need} \leq \text{Available}$

If no such i exists, go to Step(4).

Step(3):

$\text{Available} = \text{Available} + \text{Allocation}_i$
 $\text{Finish}[i] = \text{TRUE}$
 Go to Step(2).

Step(4): If $\text{Finish}[i] = \text{TRUE}$ & i :

then system is in a safe state.

* Resource Request Algorithm:
 If $\text{Request}[i][j] = k$ then process P_i wants k instances of resource type R_j .

Note: If $\text{Request}_i \leq \text{Need}_i$ go to Step(2); otherwise there exist condition because it's more than $\text{Max}[i][j]$.
 Since resources are not available.

If $\text{Request}_i \leq \text{Available}_j$ go to Step(3); otherwise P_i must wait for the old resources requested.

(3) Method to allocate requested resources to P_i by Hedge: Resources are allocated to P_i and $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$

$\text{Need}_i = \text{Need}_i - \text{Request}_i$

Step(1): Let Available & Finish be vectors of length m & n respectively. Initialize:
 $\text{Finish}[i] = \text{False} \quad \forall i=0, 1, 2, \dots, n-1$.

Hence, safe states $\prec P_1, P_3, P_4, P_2, P_0 \succ$.

Ex:- Consider a system with five processes P_0 through P_4 + three resource type A,B,C . Resource type A has 10 instances , Resource B has 5 instances & resource type C has 7 instances

Suppose that the time To , the following snapshot of system has been taken:

Process	Allocation			Max			Available			Need = Max - Allocation		
	A	B	C	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	4	5	3	3	3	2	7	4	3
P ₁	2	0	0	3	2	2	5	3	2	1	2	2
P ₂	3	0	2	9	0	2	7	4	3	6	0	0
P ₃	2	1	1	2	2	2	7	4	5	0	1	1
P ₄	0	0	2	4	3	3	10	4	7	4	3	1

check safety state

$$\text{Available} = 3 \ 3 \ 2$$

(1) Need \leq Available
 $\text{Available}_i = \text{Available} + \text{Allocation}_i$

$$\begin{aligned} \text{Available}_1 &= 3 \ 3 \ 2 + 2 \ 0 \ 0 \\ &= 5 \ 3 \ 2 \end{aligned}$$

$$\text{Available}_2 = 5 \ 3 \ 2 + 2 \ 1 \ 1 = 7 \ 4 \ 3$$

$$\text{Available}_3 = 7 \ 4 \ 3 + 0 \ 0 \ 2 = 7 \ 4 \ 5$$

$$\text{Available}_4 = 7 \ 4 \ 5 + 3 \ 0 \ 2 = 10 \ 4 \ 7$$

$$\text{Available}_0 = 10 \ 4 \ 7 + 0 \ 1 \ 0 = 10 \ 5 \ 7$$

③ \Rightarrow Update state.

Process	Allocation (A, B, C)	Need (A, B, C)	Available (A, B, C)
P ₀	(0, 1, 0)	(7, 4, 3)	(2, 3, 0)
P ₁	(2, 0, 0)	(0, 2, 0)	
P ₂	(3, 0, 2)	(6, 0, 0)	
P ₃	(2, 1, 1)	(0, 1, 1)	
P ₄	(0, 0, 2)	(4, 3, 1)	

③ check Safety Algorithm to verify if the system is in a safe state .

Step(3): Request by $P_0 = (0, 2, 0)$

① check feasibility: Request \leq Available

$$(0, 2, 0) \leq (2, 3, 0)$$

- ② Start with Available = $(2, 3, 0)$
 (b) find process P_i such that $\text{Need}_i \leq \text{Available}$

$$\rightarrow P_1 : (0, 2, 0) \leq (2, 3, 0). \text{ safe, release resources:}$$

$$\rightarrow P_3 : (0, 1, 1) \leq (2, 3, 2). \text{ safe, release resources:}$$

$$\text{Available} = (2, 1, 1) + (2, 3, 2) = (4, 4, 3)$$

$$\rightarrow P_4 : (4, 3, 1) \leq (7, 4, 3). \text{ safe, release resources:}$$

$$\text{Available} = (0, 0, 2) + (7, 4, 3) = (7, 4, 5)$$

$$\rightarrow P_2 : (6, 0, 0) \leq (7, 4, 5). \text{ safe, release resources}$$

$$\text{Available} = (3, 0, 2) + (7, 4, 5) = (10, 4, 7)$$

$$\rightarrow P_0 : (7, 4, 3) \leq (10, 4, 7) \text{ - safe.}$$

Hence Safe sequence $< P_1, P_3, P_4, P_2, P_0 >$ satisfy the safety condition. The system is safe.

Step(2): Request by $P_4 = (3, 3, 0)$

① check feasibility: Request \leq Available

$$(3, 3, 0) \leq (2, 3, 0) \text{ false}$$

Request is not granted.

Update matrices

Process	Allocation(A, B, C)	Need(A, B, C)	Available(A, B, C)
P_0	(0, 3, 0)	(7, 2, 3)	(2, 1, 0)
P_1	(3, 0, 2)	(0, 2, 0)	
P_2	(3, 0, 2)	(6, 0, 0)	
P_3	(2, 1, 1)	(0, 1, 1)	
P_4	(0, 0, 2)	(4, 3, 1)	

② check safety Algorithm:

→ Start with Available = $(2, 1, 0)$

→ Find P_i such that $\text{Need}_i \leq \text{Available}$. No process satisfies

Hence, The system is unsafe after granting the request.

Hence, the request by P_0 cannot be granted.

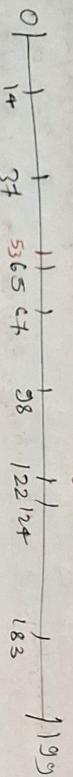
- * Default head moves to the left towards previous if previous is provided then moves towards right. (as based on previous)

(3) Disk scheduling

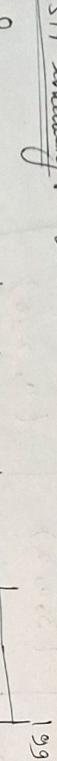
F-CFS scheduling:

FIFO manner worked in sequential order.

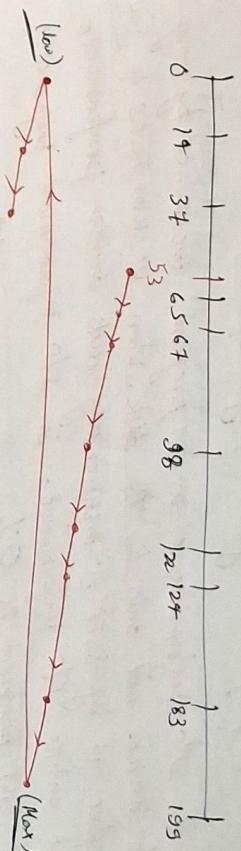
Ex- Request queue (0-199) : 58, 183, 37, 122, 14, 124, 65, 67
consider now the head pointer is in cylinder 53.



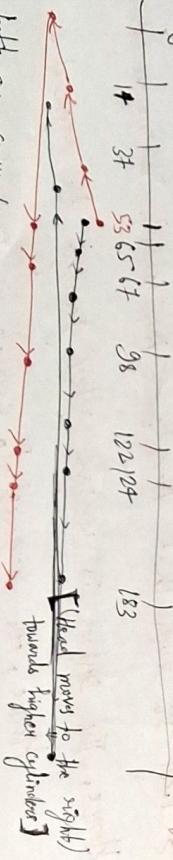
SSTF scheduling: Shortest-seek-time-first



* Look : Similar to SCAN, but the head does not go to the extreme ends (highest or lowest cylinder). It only goes as far as the farthest request in each direction before reversing.



* SCAN scheduling: The disk arm starts at one end of the disk, & moves toward the other end, servicing requests as it reaches each cylinder, until it gets to the other end of the disk. Head reversal happens after the last cylinder.



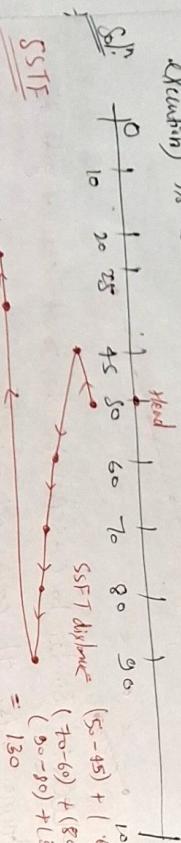
both are correct.

Ques. Consider an operating system capable of loading & executing a single program at a time. The disk head scheduling algorithm used is FCFS. If FCFS is replaced by SSTF, claimed by the vendor to give better benchmark results, what is the expected improvement in the I/O performance of user programs?

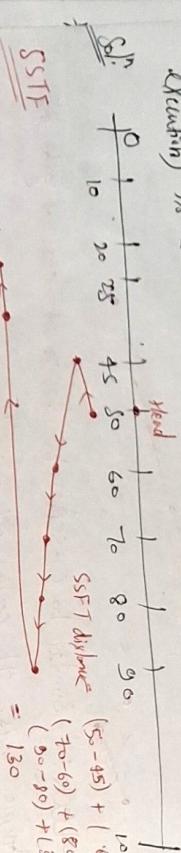
Sol. Ans. The I/O performance of a user program determined by many I/O devices not only by the disk. When we use FCFS & replace it by SSTF it improves only disk drive performance not entire I/O performance. So I/O improvement performance of user program is 0%.

Ques. Suppose the following disk request sequence (track number) for 2015-16 Q1. A disk with 100 tracks is given 45, 20, 90, 19, 50, 60, 80, 25, 70. Assume that the initial position of the R/W head is on track 50. Answer the additional distance that will be travelled by the R/W head when the SSTF algo. is used compared to the SCAN (Elevator) algo. (assuming that SCAN algo. moves towards 100 when it starts etc.)

Excluding Head SCAN is _____ blocks.



SSTF



SSFT

$$\therefore \text{SCAN distance} = 10 \text{ dm.}$$

C-Look
Given : The cylinders are numbered from 0 to 199
disk queue: 47, 38, 121, 191, 87, 17, 92, 10
The head is initially at cylinder no. 63, moving towards larger cylinder no. on its servicing path.

Final \Rightarrow Total No. of head movement

Sol. Total head movement = $(10 - 63) + (92 - 87) + (12 - 92) + (191 - 121) + (191 - 10) + (47 - 17)$
 $= 24 + 5 + 29 + 70 + 181 + 11$
 $= 346 \text{ Ans.}$

Ques. Consider the following 2D array in the C programming language. Which is stored in row-major order: int D[128][128];

Demand paging is used for allocating memory & each physical page frame holds 512 elements of the array D. The LRU page replacement policy is used by the operating system. A total of 30 physical page frames are allocated to a process which executes the following code snippet:

(APTE 2012)

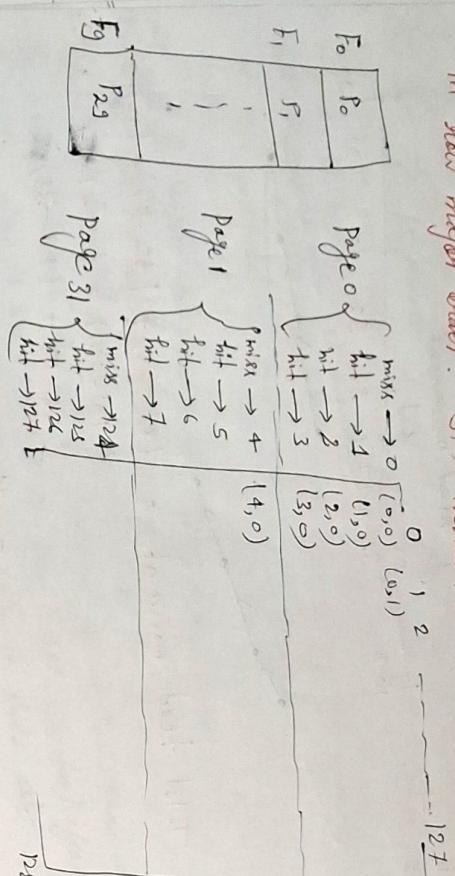
Consider the 2 functions shown below!

```

    incr() {
        decr() {
            wait(s);
            X = X + 1;
            signal(s);
        }
    }

```

Sol: Given array dimension is $D[128][128]$ which is stored in new major order. This means that $D[8]$ starts of 128 columns. The no. of page faults generated during the execution of this code snippet is $\boxed{D[5][5][5] * 10}$.



Given that, each physical frame holds 512 elements of the array D .

So page frames are allocated to a process. Array access is in column major order $[D[5][5][5] * 10]$. $\frac{512}{128} = 4$ elements

At each physical frame holds $\boxed{512}$ elements, the $\boxed{4}$ storage of the array belongs one page, we can see that, one page fault 4 elements.

$$\frac{128}{32} = 32$$

Array contains $\boxed{32}$ pages, but memory contains 30 page frames

$$\therefore \text{Total page faults} = \frac{128 \times 128}{4} = \frac{2^{14}}{2^2} = \boxed{4096}$$

Ques.

Sol: Let V_1 & V_2 be the values of X at the end of execution of all the threads with implementations I_1 & I_2 respectively, which one of the following choices corresponds to the minimum possible values of V_1 & V_2 respectively?

- (A) 15, 7 (B) 7, 7 (C) 12, 4 (D) 12, 8.

Sol: I_1 : s is a binary semaphore initialized to 1. With binary semaphores there can't be 2 threads that can simultaneously access the critical section. Thus, each thread will enter $incr()$ or $decr()$ one by one. Hence, only possible value of X at the end is 12.

T2: With counting semaphore of value 2. Two processes at a time can access critical section.

$$X = X + 1$$

↓

Load X
 Add #1, X
 Store X

$$X = X - 1$$

↓

Load X
 Sub #1, X
 Store X

In order to minimize X value, we need to maximize the impact of devr() of devl() to minimize the impact of incr().

devr()
devl()

decr()

Load X

Held
X=7
Run incr() 5 times

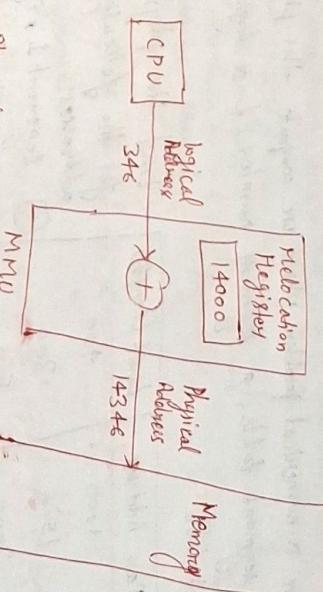
Sub #1, X
Store X

Run last part.

$$X = 7$$

Final.

Logical vs Physical Space



→

An address generated by the CPU is commonly referred to as a logical address while a virtual address is an address seen by the main memory unit (MMU) is commonly referred to as physical address (PA).

→ The set of all logical addresses generated by a program is a LVA space and the set of all physical addresses is a PVA space.

T2: Initially, $X = 10$, $\lambda = 1$ (mutual exclusion) of ix binary semaphore. So after executing 3 incr(), $X = 10 + 5 = 15$
if after executing 3 devr(), $X = 15 - 3 = 12$

i.e. 2 processes can enter the CS at the same time.

15 \neq 7

i.e. correct answer,

12 7

① Necessary condition of Deadlock

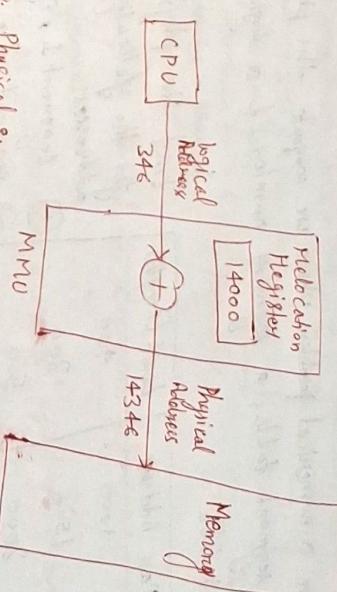
Mutual exclusion
No preemption
Circular wait

Deadlock avoidance
Recover Rollback
Recover Process Termination

Method to recovery from

② How logical address to physical address mapped. Explain in details.

Ans:

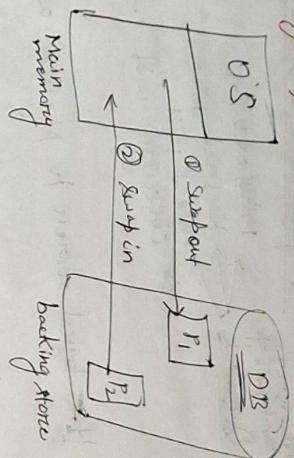


→ logical & physical addresses are the same in compile-time & load-time address binding scheme; logical (virtual) & physical (actual) addresses differ in execution-time address binding scheme.

The Memory Management Unit is a new device that maps virtual to physical address. In MMU scheme, the value in the relocation register is added to every address generated by a user program.

at time it is sent to memory as follows:

- * ① Dynamic Loading
- * ② Dynamic Linking
- + ③ Overlays
- * ④ Swapping → Roll Out



- (3) Let's solve a numerical problem where we compute the physical address given a segment table, a segment number & a logical address.

Given segment table

Seg. No.	Base Address	Limit (bytes)
0	1000	400
1	1500	300
2	2000	500
3	3000	250

Find the physical address for the following logical addresses:

- ① Segment 1, offset 120
- ② Segment 2, offset 450
- ③ Segment 3, offset 300

Soln
① check if the offset is within the limit of the segment.
→ If offset \leq limit, then logical address is valid.
Otherwise, proceed to the next step.

- ② Compute the physical address.

$$PA = \text{Base Address of Segment} + \text{Offset}$$

$$\text{① Segment 1, offset 120 : } \begin{cases} 120 \leq 400 & (\text{Valid}) \\ 120 + 1000 = 1120 & \end{cases}$$

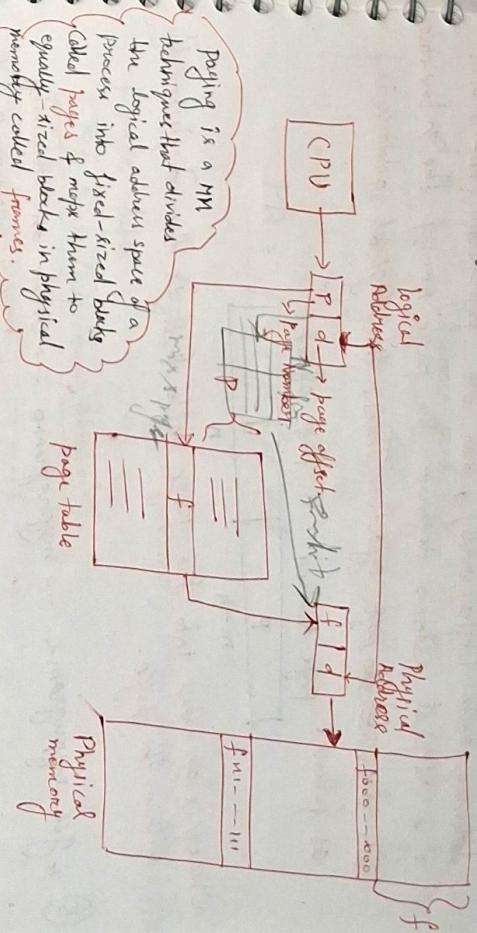
$$\text{② Segment 2, offset 450 : } \begin{cases} 450 \leq 500 & (\text{Valid}) \\ 450 + 2000 = 2450 & \end{cases}$$

$$\text{③ Segment 3, offset 300 : } \begin{cases} 300 \leq 250 & (\text{Invalid}) \\ 300 + 3000 = 3300 & \end{cases}$$

Hence logical address is invalid.

④ How paging is mapped from logical address to physical address

→ Eliminate external fragmentation.



→ Main memory is divided into a number of equal-size blocks, called frames.

→ Each process is divided into a no. of equal-size blocks, called pages.

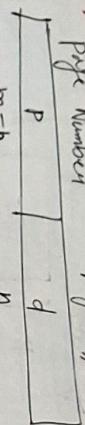
→ A process is loaded by loading all of its pages into available frames (may not be contiguous).

→ Process of translation from logical to physical address.

→ Every address generated by the CPU is divided into 2 parts → Page Number (P) & page offset (L). The page number is used as an index into a page table.

→ The page table contains the base address of each page in physical memory. This base address is combined with the page offset to define the physical address that is sent to the memory unit.

↳ If the size of logical address space is 2^m & a page size is 2^n
 ; a disuniting units (blocks) words), then the high-order ($m-n$) bits
 of a logical address designate the page no. & the low-order
 bits designate the page offset. Thus, the logical address is
 as follows: Page Number Page offset.



Where p is an index into the page table & d is the displacement
 within the page.

(5) * $\text{gcc } -f \text{ myprogram.c} \rightarrow \text{myprogram.o}$ → Create an executable
 named myprogram.o. & file type Executable.

* $\text{gcc } -f \text{ myprogram.c} \rightarrow \text{myprogram.o} \rightarrow \text{create an object file}$
 named myprogram.o.
 ↳ flag to generate an object file.
 & file type → object file.

* How paging works:

= ① Logical address \rightarrow Page No. (P) $\quad \exists L A = P + d = \text{Page No.} + \text{offset.}$

② Physical address \rightarrow Frame No. (f) $\quad \exists PA = \text{Frame No.} \times \text{Page size} + \text{offset.}$

Ex- Page size = $4KB = 2^{12}$ Bytes
 then $P = 32 - 12 = 20$ bits (page Number)
 $d = 12$ (offset).

* Mapping logical address to physical Address

↳ CPU generates Logical address
 ↳ Decompose logical address $\rightarrow P \# d$

$$p = \lfloor \text{Logical Address} \rfloor \text{ mod Page size.}$$

↳ Access the Page table : find corresponding frame for p .
 ↳ Calculate physical address:

$$PA = f * \text{Page size} + d$$

Page Number (P)	Frame Number (F)
0	5
1	3
2	4

Sol: ① Decompose logical address:

$$\begin{aligned} p &= \lfloor 1200 / 256 \rfloor = 4 \\ d &= 1200 \bmod 256 = 176 \end{aligned}$$

② Access Page table : Page 4 is not mapped in the given page table. Invalid logical address.

↳ Logical Address = 100 then
 $\exists p = \lfloor 100 / 256 \rfloor = 1 \quad \& \quad d = 100 \bmod 256 = 144$

- ① $p = 1$ maps to frame 3.
- ② $PA = 3 * 256 + 144 = 912$.

Hence page table is key to translating logical address into PA.
 ↳ Decompose logical address $\rightarrow P \# d$

④ Find the no. of child processes given the pseudo code with fork.

From $(i=0; i \geq 10; i++)$

Hence, No. of child processes = $2^n - 1 = 2^5 - 1$

if $i/2 = 0:$

fork();

= (3) Ans.

⑧ Finding the best algorithm b/w FCFS-fit & Best-fit given the size of five memory blocks of size of 4 processes.

Given: Block size = {100, 500, 200, 300, 600}

Process size = {212, 417, 112, 426}

Sol: ① First-Fit Alg: Allocate the process to the first block that is large enough to accommodate it.

Process	Block size	Block Allocated	Remaining Block Size After Allocation
P ₁	212	Block 2 (500)	500 - 212 = 288
P ₂	417	Block 5 (600)	600 - 417 = 183
P ₃	112	Block 3 (200)	200 - 112 = 88
P ₄	426	No Block	Not Allocated.

P₁, P₂, P₃ are allocated but P₄ is not allocated.

Best-fit Alg: Allocate the process to the smallest block.

That is large enough to accommodate it.

Fragmentation elimination techniques

① Paging → Eliminates external fragmentation (Internal fragmentation may be reduced).

② Segmentation → Reduces internal fragmentation but can still lead to external fragmentation.

③ Compaction → Rearranges memory to group all free spaces together. Reduces external fragmentation but is time-consuming.

④ Buddy System → Balances internal fragmentation by combining free blocks into powers of two.

Q. Is there any fragmentation possible depending on the fixed size of block? When size of block is not fixed, is there possibility of no fragmentation? Give explanation with reason.

Ans: Fragmentation with fixed Block Sizes → Internal fragmentation Static allocation

Ex: → Block size = 100 kB → Process size = 70 kB

→ Allocated block = 100 kB → Wasted memory = 100 - 70 = 30 kB.

Thus, fixed-size blocks always have the potential for a fragmentation because processes rarely fit perfectly into a block.

* Fragmentation with Variable Block Sizes → dynamic allocation External fragmentation.

Ex: → Initially, memory has 100 kB, 300 kB, 200 kB free blocks. → Process 1 (120 kB) is allocated to the 300 kB block, leaving 300 - 120 = 180 kB as free block.

Now, there's 100 kB, 180 kB, & 200 kB free.

→ Process 2 (250 kB) cannot be allocated even though the total free memory is 100 + 180 + 200 = 480 kB.

Q10 Consider a system running ten I/O-bound tasks and one CPU-bound task. Assume that the I/O-bound task issue an I/O operation once for every millisecond of CPU computing so that each I/O operation takes 10 milliseconds to complete. Also assume that the context-switching overhead is 0.1 millisecond and that all processes are long-running tasks. Describe the CPU utilization for a Round-Robin scheduler when:

- (a) The time quantum is 1 ms. (b) The time quantum is 10 ms.
- What may be the burst time of CPU-bound process (but it should be greater than TQ), it will run upto 4 TQ, often that already I/O bound process should come into Ready Queue + Total time utilized = $11 * 10^3 + 1$ ten times of I/O bound process + 4 TQ for CPU-bound.

→ Extra time consumed = 1 context-switch for each process = $10 * 0.1$

$$\rightarrow \text{Extra time consumed} = 1 \text{ context-switch for each process} = 10 \text{ ms.}$$

$$\text{① If } TQ = 1 \text{ ms} \Rightarrow \text{Efficiency} = \frac{10 \text{ ms} + 1 \text{ ms}}{10 \text{ ms} + 1 \text{ ms} + 1.1 \text{ ms}} = \frac{11}{12.1} = 90.9\%$$

$$\text{② If } TQ = 10 \text{ ms} \Rightarrow \text{Efficiency} = \frac{10 \text{ ms} + 1 \text{ ms}}{10 \text{ ms} + 10 \text{ ms} + 1.1 \text{ ms}} = \frac{20}{21.1} = 94.78\%$$

⑥ Process synchronization of mutex related pseudo code-

→ A situation where several processes access manipulate the same data concurrently & the outcome of the execution depends on the particular order in which the access takes place, is called a race condition.

Producer - Consumer Problem

- Unbounded buffer places no practical limit on the size of the buffer.
- Bounded buffer assumes that there is a fixed buffer size.

* Bounded Buffer - shared Memory Solution

→ Used to illustrate the power of synchronization primitives.
→ we assume that the pool consists of n buffers, each capable of holding one item. The mutex semaphore provided mutual exclusion for access to the buffer pool & is initialized to 1.
→ The empty & full semaphore count the no. of empty & full buffers respectively.
→ The semaphore empty is initialized to the value n; the semaphore full is initialized to the value 0.

The code for the producer process is

do {
 produce an item in nextOp
 Wait (empty);
 Wait (mutex);

remove an item from buffer to
nextOp
Signal (mutex);
Signal (empty);
Consume the item in nextOp
} while (1);

signal (mutex);
signal (full);
if while (1);

signal (mutex);
signal (empty);
Consume the item in nextOp
} while (1);

a race condition.

The Reader-Writers Problem

→ The semaphore number of `wint` are initialized to 1;
→ `readcount` is initialized to 0.

The code for a writer process is

```
do {
    wait (wint);
    writing is performed
    signal (wint);
    while (!);
}
```

Note: if n writers are waiting, then one reader is

queued on `wint`, & n-1 readers are queued on `readcount`. Also observe that when a writer executes `signal (wint)`, we may resume the execution of either `while (!);` the waiting readers or a different waiting writer.

```
wait (mmtx);
readcount --;
if (readcount == 0)
    signal (wint);
    signal (mmtx);
reading is performed
wait (mmtx);
```

The Dining-Philosophers Problem

The structure of philosophers: & the semaphore chopstick [5]; where do I wait (chopstick[i]); & where all the elements of wait (chopstick[(i+1) % 5]), eat

```
signal (chopstick[i]);
signal (chopstick[(i+1) % 5]);
```

think
} while (!);
} while (!);

- binary semaphore S_1 & S_2 ;
- int c;
- Initially $S_1 = 1$, $S_2 = 0$ & the value of integer C is set to the initial value of the counting semaphore S.
- * Wait operation on the counting semaphore S:
- Wait (S_1);
- C--;
- if ($c < 0$) {
- if ($c \leftarrow 0$) {
- signal (S_2);
- } else
- signal (S_1);

* Signal operation on the S:

Wait (S_1);

$C++;$

if ($c < 0$) {

 signal (S_2);

 wait (S_2);

} signal (S_1);

↓

↓

↓

↓

Peterson's solution: is a SW based sol'n to the critical section problem.

Consider 2 processes P_0 & P_1 . For convenience, when presenting P_i ,

we use P_i to denote the other process that is $j = 1-i$.

→ The process sharing 2 variables:

boolean flag[2];

int turn;

Initially $\text{flag}[0] = \text{flag}[1] = \text{false}$. If value of `turn` is increased

↓ (but either 0 or 1).

The structure of process P_i is shown below:

(1) Mutual exclusion is preserved.

(2) The progress requirement is satisfied.

(3) The bounded-waiting requirement is met.

do {

 flag[i] = true

 turn = j

 while (flag[j] & turn == j):

 Critical section;

 flag[i] = false

};

lime (1);

lime (1);

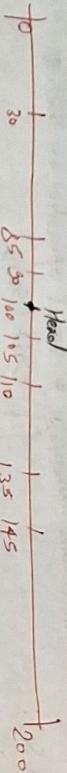
Comprehensive June 2021 QP with Solutions

① 0, 430 :

$$PA = 1219 + 430 = 1649$$

- ⑤ Given : disk = 801 cylinders, Head position = 100
 Queue = 30, 85, 90, 100, 105, 110, 135, 145

SSTF algorithm to find sequence in which these requests are processed / serviced.



$$= (100 - 105) + (110 - 105) + (110 - 90) + (90 - 85) + (135 - 85) +$$

$$(145 - 135)$$

$$= 5 + 5 + 20 + 5 + 50 + 115 + 10$$

210 char

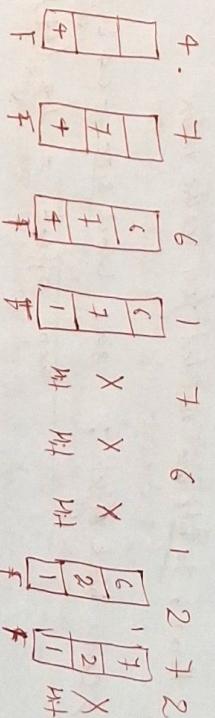
Given Segment table

Seg No.	Base	Length
0	1219	700
1	2300	11
2	30	100
3	1327	580
4	1952	36

which of the following logical addresses will produce trap addressing errors?

$$\begin{array}{l} \textcircled{1} 0, 430 \\ \textcircled{2} 1, 11 \\ \textcircled{3} 2, 100 \\ \textcircled{4} 3, 425 \\ \textcircled{5} 4, 95 \end{array}$$

$$\begin{aligned} \text{Soln} &: \\ &\text{Page reference string : } 4, 7, 6, 1, 7, 6, 1, 2, 7, 2 \\ &\text{Page frame } = 3 \quad (\text{given}), \text{ FIFO } \& \text{ find total no. of page faults} \\ &\text{hit ratio } \& \text{ miss ratio.} \end{aligned}$$



$$\begin{aligned} \text{Total Page faults} &= 6 \\ \text{Page Hits} &= 10 - 6 = 4 \\ \text{Hit Ratio} &= \frac{\text{Page Hits}}{\text{Total References}} = \frac{4}{10} = 0.4 \text{ or } 40\% \end{aligned}$$

- Soln:
 ① Valid offset check : if $\text{offset} < \text{length}$ (valid) else invalid.
 ② $\text{offset} > \text{seg.length} \Rightarrow$ trap addressing error.

- ② PA = Base Address + offset (if valid).

⑥ Given:

Virtual address size \rightarrow 46 bits
Physical address size = 32 bits

Size of page = ?
Page table organization = 3-level paging.
PTE size = 32 bits (4 bytes)

* Cache properties : Direct cache size : $1\text{MB} \approx 2^{20}$ bytes

$$\begin{aligned} \text{Set associativity} &= 16\text{-way} \\ \text{Block size} &= 64 \text{ bytes} \end{aligned}$$

* Cache indexing : Virtually indexed, physically tagged.

$$\text{S!} \quad \text{No. of refs} = \frac{\text{cache size}}{\text{associativity} \times \text{block size}} = \frac{2^{20}}{16 \times 64} = 2^{10} = 1024 \text{ refs.}$$

Hence, 10 bits are needed for cache index.

\rightarrow Cache block offset : Since block size = 64 bytes
Then offset = $\log_2(64) = 6$ bits.

$$\rightarrow \text{Page offset} = \text{Block offset} + \text{Cache index} = 6 + 10 = 16 \text{ bits}$$

$$\rightarrow \text{Hence, Page size} = 2^{16} \text{ bytes} = 64 \text{ KB. size}$$

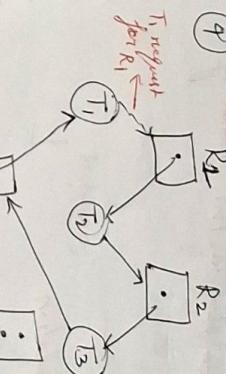
Given : $\forall i, R_i \in \{R_1, R_2, R_3, R_4\}$

Given : T_1 holds R_3 & request R_1 (dotted arrow).

Given : T_2 holds R_1 & request R_2 .

Given : T_3 holds R_2 & request R_3 .

④



To determine if a deadlock occurs, we analyze the scenario based on the 4 necessary & sufficient conditions for deadlock.

Given : $T_1 \rightarrow T_2 \rightarrow \text{read-only}$
 $T_3 \rightarrow \text{write-only}$
 $T_1, T_2, T_3 \rightarrow R_1, R_2, R_3, R_4$.

② Existing Allocation : $\rightarrow T_1$ holds R_3 & request R_1 (dotted arrow)

$\rightarrow T_2$ holds R_1 & request R_2 .

$\rightarrow T_3$ holds R_2 & request R_3 .

③ If T_1 requests R_1 : A circular dependency is created.

$\rightarrow T_1 \rightarrow R_1 \Rightarrow$ Requested by T_1 & held by T_2

$\rightarrow T_2 \rightarrow R_2 \Rightarrow$ Requested by T_2 & held by T_3

$\rightarrow T_3 \rightarrow R_3 \Rightarrow$ Requested by T_3 & held by T_1 .

Deadlock Analysis

① Mutual Exclusion : Yes, resources R_1, R_2, R_3, R_4 are assigned to one process at a time. Thus condition is satisfied.

② Hold & Wait : Yes, condition satisfied.

③ No preemption : Condition satisfied.

$$\begin{aligned} \text{Tag?} \quad \text{TLB} &= 128 \text{ page table entries of } 4\text{-way set associative.} \\ \text{No. of sets} &= \frac{\text{Total TLB entries}}{\text{Associativity}} = \frac{128}{4} = 32 \text{ sets.} \\ \text{# cache index} &= \log_2(32) = 5 \text{ bits.} \end{aligned}$$

True

True

True

True

②

int function (int level) {

 int i = 0;
 for (i = 0; i < level; i++) {

 if (fork() && fork())
 break;

}

 P₀ & P₁ both proceed to execute
 the second fork();

 P₁ creates a new process (P₂).

 Total Process = P₀, P₁, P₂, P₃.

③

After the 2nd fork():

 only the parent processes (P₀ & P₁) satisfy the
 of (fork() && fork()) condition.

 These parent processes execute the break statement, exiting the loop.

 Subsequent iterations (i = 1, 2, ..., level - 1):

 The break statement ensures that no further iterations of the loop
 are executed.

 Therefore, no additional processes are created after the
 first iteration.

Total No. of Processes:

In 1st iteration, the fork() calls result in $2^0 = 1$ process.
For the entire program:

 Initial Process (P₀) = 1 fork()

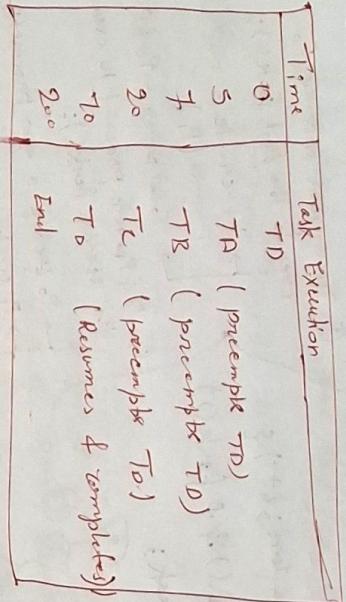
 Additional processes created in the first iteration = 3 new forks

 Total No. of processes is = $2^{min(2, level)} = 2^2 = 4$

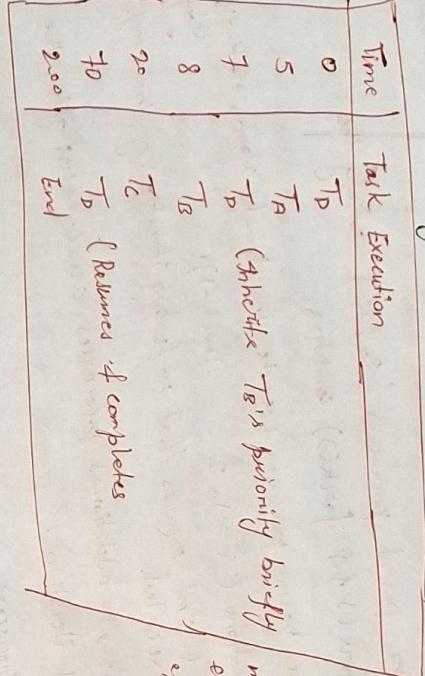
Ans

- ③ Threads share the heap but not the stack. What would happen if the threads used the same stack for execution? If threads used the same stack for execution, several critical problems would arise due to the nature of stack memory.
- Data corruption
 - Loss of thread isolation
 - Incorrect fund'n behavior
 - Race conditions
 - Stack overflow
- Why threads have separate stacks
- Each Thread is given its own stack so it can independently manage its local execution context (func'ns, local variables, return addresses).
- Threads share the heap for global or dynamically allocated data, which allows them to communicate & share stack.

① ② Without priority inheritance



③ With priority inheritance



* Priority inversion might cause delays.

* So MCQ + Descriptive questions.

* Related of OS basics & GATE MCQ.
→ Simple Linux commands.

* 1 Question related C-Lock

* 2 Question related Memory management

* 2 Question Deadlock, conditions & How to prevent from deadlock methods/techniques explanation with example.

* 1 Question Dining philosopher & its effectiveness.

* 1 Context switching related.

* 1 Question LRU related in partitions.

* Priority inheritance mitigates the issue, ensuring smoother execution.

(a) Using priority

(b) Using priority